

An Intro to Distributed Memory Computing and MPI

Jack Poulson, Rob Schreiber
Stanford ICME

Class 1

Parallel computing

Moore's Law

Dennard Scaling

The end of scaling

Scientific computing == Parallel computing

Big data sets

Swapping and virtual memory

Economics of big memory


(Sequoia machine -- 1,600,000 gigabytes of memory)

MIMD, SIMD, SPMD (models of) parallel machines


- SIMD: All the processors share one program and program counter, but each processor has its own data
- MIMD: Each processor runs its own program, has its own program counter
- SPMD: Like MIMD, but you write one program and a copy runs on each processor

Computer architectures

A single core computer




A single processor vector/SIMD computer




Two kinds of (MIMD) parallel machines

- Shared-memory multiprocessors
 - “Scale Up”
 - Hard to build, easy to program
- Distributed-memory multiprocessors
 - A.K.A. Multicomputer
 - “Scale Out”
 - Easy to build, hard to program


Shared memory multiprocessor




Shared memory multiprocessor (really)



An old distributed memory multicomputer: 1980-2000



An newer distributed memory multicomputer



Supercomputers Today

- They are (essentially) all multicomputers with complex, multi-core nodes



Supercomputers in 2013

- 10^{16} ops/sec; 10^{15} bytes of memory
- 10,000s of square feet
- Hundreds of racks
- Fiber optic network
- Water cooled
- Cost \$100M and up
- Used for the biggest problems
- **Nearly all code is MPI code**

Parallel Programming

- Take sequential programs; compile them into parallel executables!
 - No work!
 - No performance ☹️
- Write programs that are sequential, but implicitly parallel

Implicitly parallel programs

```
REAL *8 A(1000), B(1000), C(1000)
```

```
C = A + B
```

(Compiler vectorization)

```
DO I = 1, 1000
```

```
    C(I) = A(I) + B(I)
```

Shared memory, threads

```
cilk int fib (int n) {  
    if (n < 2) return n;  
    else {  
        int x, y;  
        x = spawn fib (n-1);  
        y = spawn fib (n-2);  
        sync;  
        return (x+y);  
    }  
}
```


How do you write and run SPMD programs?

- On your PC (phone, phablet):
 - `gcc myprog.c -o node.out`
- `Bigcluster_login>> mpirun -np 512 node.out`
- Starts `node.out` on each of 512 nodes of the cluster.

Embarrassingly parallel problems


- Every node computes something, independent of the other nodes
- They share nothing
- They do not communicate
- Efficient if it is load-balanced

Can the nodes actually cooperate? Share data?

- There is no shared memory with which to share data
- How can they cooperate?
- Send messages
- Like the postal service: message goes in at one node, with a destination address, the communication system delivers it (reliably, later) to the intended recipient

A message is sent and received

- Node A: `send(A_buffer, len, node_B_id)`
- Node B: `recv(B_buffer, len, node_A_id)`
- Messaging library hides LOTS of details → portability



MPI -- hello world

```
#include "mpi.h"
int main ( int argc, char *argv[] )
{
    int me; int ierr; int nprocs; double wtime;
    ierr = MPI_Init ( &argc, &argv );
    ierr = MPI_Comm_size ( MPI_COMM_WORLD, & nprocs );
    ierr = MPI_Comm_rank ( MPI_COMM_WORLD, &me);

    /* p0 is the spokesprocess */
    if ( me == 0 ) {
        wtime = MPI_Wtime ( );
        printf ( " The number of processes is %d.\n", nprocs );
    }
    printf ( " Process %d says 'Hello, world!'\n", me);
    if ( id == 0 ) {
        printf ( "That took = %f seconds.\n", MPI_Wtime ( ) - wtime);
    }
    ierr = MPI_Finalize ( );
}
```

MPI -- communication

```
{
    int ierr, me, nprocs, other_rank, sum_of_ranks, items, tag;
    MPI_Status status;
    ierr = MPI_Init ( &argc, &argv );
    ierr = MPI_Comm_size ( MPI_COMM_WORLD, & nprocs );
    ierr = MPI_Comm_rank ( MPI_COMM_WORLD, &me);

    /* p0 adds the ranks*/
    if ( me == 0 ) {
        sum_of_ranks = 0; items = 1; tag = 0;
        for (int i = 1; i < nprocs; i++) {
            ierr = MPI_Recv(&other_rank, items, MPI_INT, i, tag, MPI_COMM_WORLD,
                &status);
            sum_of_ranks += other_rank;
        }
        printf ( " All together, our ranks total %d \n", sum_of_ranks );
    } else {
        MPI_Send ( &me, items, MPI_INT, 0, tag, MPI_COMM_WORLD );
    }
    ierr = MPI_Finalize ( );
}
```

Programming supercomputers

- Multicomputers -- nothing is shared
- A process can see only its own data
- The processes, one per processor, send and receive messages in order to cooperate
- MPI is the standard way to do that

MPI -- Six basic functions

- Init, Finalize, Comm_size, Comm_rank, Send, Recv
- You are now ready to write some MPI programs!