

High-performance computation of pseudospectra

Jack Poulson^{1,2,3} Greg Henry⁴

¹Dept. of Mathematics
Stanford University

²Inst. of Comput. and Math. Eng.
Stanford University

³School of Comput. Science and Eng.
Georgia Institute of Technology

⁴Intel Corporation

SIAM CSE15
Salt Lake City, Utah
March 14, 2015

Context / Trojan horse for current research

- ▶ Have spent the past few months developing distributed sparse-direct Interior Point Methods for QPs,

$$\begin{aligned} \min_x \quad & \frac{1}{2}x^T Qx + c^T x \\ \text{s.t.} \quad & Ax = b, Gx + s = h, s \geq 0 \end{aligned}$$

- ▶ Ran into fundamental instability of dynamic regularization [Altman/Gondzio-1998, Li/Demmel-1998] closely related to pseudospectra
- ▶ There exists a simple, small *symmetric quasi-definite matrix* [Vanderbei-1993] with condition number less than 15 that can cause dynamic regularization to catastrophically fail

Context / Trojan horse for current research

- ▶ Have spent the past few months developing distributed sparse-direct Interior Point Methods for QPs,

$$\begin{aligned} \min_x \quad & \frac{1}{2}x^T Qx + c^T x \\ \text{s.t.} \quad & Ax = b, Gx + s = h, s \geq 0 \end{aligned}$$

- ▶ Ran into fundamental instability of dynamic regularization [Altman/Gondzio-1998,Li/Demmel-1998] closely related to pseudospectra
- ▶ There exists a simple, small *symmetric quasi-definite matrix* [Vanderbei-1993] with condition number less than 15 that can cause dynamic regularization to catastrophically fail

Context / Trojan horse for current research

- ▶ Have spent the past few months developing distributed sparse-direct Interior Point Methods for QPs,


$$\begin{aligned} \min_x \quad & \frac{1}{2}x^T Qx + c^T x \\ \text{s.t.} \quad & Ax = b, Gx + s = h, s \geq 0 \end{aligned}$$

- ▶ Ran into fundamental instability of dynamic regularization [Altman/Gondzio-1998,Li/Demmel-1998] closely related to pseudospectra
- ▶ There exists a simple, small *symmetric quasi-definite matrix* [Vanderbei-1993] with condition number less than 15 that can cause dynamic regularization to catastrophically fail

The irrelevance of pivot magnitudes

$$M_n \equiv \begin{pmatrix} 1 & 2 & 2 & & & \\ 2 & 5 & 2 & & & \\ & 2 & 5 & \ddots & & \\ & & \ddots & \ddots & & \\ & & & & 5 & 2 \\ & & & & 2 & 5 \end{pmatrix} = \begin{pmatrix} 1 & & & & & \\ 2 & 1 & & & & \\ & 2 & 1 & & & \\ & & \ddots & \ddots & & \\ & & & & 1 & \\ & & & & 2 & 1 \end{pmatrix} \begin{pmatrix} 1 & 2 & 2 & & & \\ & 1 & 2 & & & \\ & & 1 & \ddots & & \\ & & & \ddots & & \\ & & & & 1 & 2 \\ & & & & 2 & 1 \end{pmatrix}$$

While all eigenvalues of $L_n \equiv \text{chol}(M_n)$ are 1, $\|(L_n - \xi I)^{-1}\|_2$ is large within a disk containing the origin



The irrelevance of pivot magnitudes

$$M_n \equiv \begin{pmatrix} 1 & 2 & & & & \\ 2 & 5 & 2 & & & \\ & 2 & 5 & \ddots & & \\ & & \ddots & \ddots & \ddots & \\ & & & & 5 & 2 \\ & & & & 2 & 5 \end{pmatrix} = \begin{pmatrix} 1 & & & & & \\ 2 & 1 & & & & \\ & 2 & 1 & & & \\ & & \ddots & \ddots & & \\ & & & & 1 & \\ & & & & 2 & 1 \end{pmatrix} \begin{pmatrix} 1 & 2 & & & & \\ & 1 & 2 & & & \\ & & 1 & \ddots & & \\ & & & 1 & \ddots & \\ & & & & 1 & 2 \\ & & & & 2 & 1 \end{pmatrix}$$

But the condition number of

$$K_n \equiv \begin{pmatrix} M_n & I \\ I & -I \end{pmatrix}$$

for $n = 100$ is less than 15 while causing dynamic regularization to fail catastrophically when solving against $L_n = \text{chol}(M_n)$.

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent norm over vertical line in complex plane [Van Loan-1985]
- ▶ Proposed as means of computing distance to nearest stable matrix; counter-example given in [Demmel-1987]
- ▶ Counter-example involved first (computer) pseudospectral plot

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent norm over vertical line in complex plane [Van Loan-1985]
- ▶ Proposed as means of computing distance to nearest stable matrix; counter-example given in [Demmel-1987]
- ▶ Counter-example involved first (computer) pseudospectral plot


Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent norm over vertical line in complex plane [Van Loan-1985]
- ▶ Proposed as means of computing distance to nearest stable matrix; counter-example given in [Demmel-1987]
- ▶ Counter-example involved first (computer) pseudospectral plot

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent norm over vertical line in complex plane [Van Loan-1985]
- ▶ Proposed as means of computing distance to nearest stable matrix; counter-example given in [Demmel-1987]
- ▶ Counter-example involved first (computer) pseudospectral plot

$$D(n, \beta) = (\beta J_{-\beta^{-1}, n})^{-1}, \quad n = 3, \beta = 100$$



Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra...*

Most common tool is EigTool [Wright-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra...*

Most common tool is EigTool [Wright-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra...*

Most common tool is EigTool [Wright-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra...*

Most common tool is EigTool [Wright-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra...*

Most common tool is EigTool [Wright-2001], but level 2 BLAS (trsv) and sequential

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \{ \xi \in \mathbb{C} : \sigma_{\min}(\xi I - G) < \epsilon \}.$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \{ \xi \in \mathbb{C} : \sigma_{\min}(\xi I - G) < \epsilon \}.$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \{ \xi \in \mathbb{C} : \sigma_{\min}(\xi I - G) < \epsilon \}.$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm for $\mathcal{L}_\epsilon^2(A)$

Algorithm: Two-norm pseudospectra via extended Van Loan algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shifts $\Omega \subset \mathbb{C}$, restart size k

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(\xi I - A)^{-1}\|_2\}_{\xi \in \Omega}$

$G := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

 // Estimate $\|(\xi I - G)^{-1}\|_2$ via Restarted Arnoldi

 Choose $v_0 \in \mathbb{C}^n$ with $\|v_0\|_2 = 1$

while not converged do

for $j = 0, \dots, k - 1$ **do**

 // $(\xi I - G)^{-H}(\xi I - G)^{-1} V_j = V_j H_j + v_j(\beta_j e_j)^H$

$x_j := (\xi I - G)^{-H}(\xi I - G)^{-1} v_j$

 Expand Arnoldi decomposition using x_j

$[\lambda, v_0] := \text{MaxEig}(H_k)$

$\phi(\xi) := \text{RealPart}(\lambda)$

[Hager-1984,Higham-1988] approach for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, $\Omega \subset \mathbb{C}$

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(A - \xi I)^{-1}\|_1\}_{\xi \in \Omega}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

 // Estimate $\|(A - \xi I)^{-1}\|_1$ via Hager-Higham algorithm

 Choose $u \in \mathbb{C}^n$ with $\|u\|_1 = 1$

$k := 0$

repeat

$x := u$

$y := Q(G - \xi I)^{-1} Q^H x$

$w := Q(G - \xi I)^{-H} Q^H \text{sign}(y)$

$u := e_j$ where $|w(j)| = \|w\|_\infty$

$k := k + 1$

until ($k \geq 2$ and $\|w\|_\infty \leq w^H x$) or $k \geq 5$

$\phi(\xi) := \|y\|_1$

 // Take the maximum between the current estimate and a heuristic

$b := \frac{2}{3n} [(-1)^j (n + j - 1)]_{j=0:n-1}$

$x := Q(G - \xi I)^{-1} Q^H b$

$\phi(\xi) := \max(\phi(\xi), \|x\|_1)$

Block algorithm [Higham/Tisseur-2000] should be used in practice

[Hager-1984,Higham-1988] approach for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, $\Omega \subset \mathbb{C}$

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(A - \xi I)^{-1}\|_1\}_{\xi \in \Omega}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

 // Estimate $\|(A - \xi I)^{-1}\|_1$ via Hager-Higham algorithm

 Choose $u \in \mathbb{C}^n$ with $\|u\|_1 = 1$

$k := 0$

repeat

$x := u$

$y := Q(G - \xi I)^{-1} Q^H x$

$w := Q(G - \xi I)^{-H} Q^H \text{sign}(y)$

$u := e_j$ where $|w(j)| = \|w\|_\infty$

$k := k + 1$

until ($k \geq 2$ and $\|w\|_\infty \leq w^H x$) or $k \geq 5$

$\phi(\xi) := \|y\|_1$

 // Take the maximum between the current estimate and a heuristic

$b := \frac{2}{3n} [(-1)^j (n + j - 1)]_{j=0:n-1}$

$x := Q(G - \xi I)^{-1} Q^H b$

$\phi(\xi) := \max(\phi(\xi), \|x\|_1)$

Block algorithm [Higham/Tisseur-2000] should be used in practice

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Generalized multishift (quasi-)triangular solves can be handled in an analogous manner (ask about appendix if interested)
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

Interleaved Van Loan algorithm for $\mathcal{L}_\epsilon^2(A)$

Algorithm: Two-norm pseudospectra via interleaved extended Van Loan algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$, restart size k

Output: $f \approx [\|(z(s)I - A)^{-1}\|_2]_{s=0:m-1}$

$G := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

Initialize each column of $W_0 \in \mathbb{C}^{n \times m}$ to have unit two-norm

$\mathcal{I} := (0, \dots, m-1)$

while $\mathcal{I} \neq \emptyset$ **do**

for $j = 0, \dots, k-1$ **do**

 // $(z(s)I - G)^{-H}(z(s)I - G)^{-1}\mathcal{W}_j(t) = \mathcal{W}_j(t)\mathcal{H}_j(t) + \mathcal{W}_j(:, t)(b_j(t)\mathbf{e}_j)^H, \forall s = \mathcal{I}(t)$

$X := \text{MultishiftSolve}(G, z(\mathcal{I}), \mathcal{W}_j)$

$X := \text{MultishiftSolve}(G^H, \bar{z}(\mathcal{I}), X)$

 Expand Arnoldi decompositions

foreach $s = \mathcal{I}(t)$ **do**

$[\lambda, W_0(:, t)] := \text{MaxEig}(\mathcal{H}_k(t))$

$f(s) := \text{RealPart}(\lambda)$

if converged **then**

 Delete $\mathcal{I}(t)$ and $W_0(:, t)$

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

Initialize each column of $U \in \mathbb{C}^{n \times m}$ to have unit one-norm

$\mathcal{I} := (0, \dots, m-1)$

$k := 0$

while $\mathcal{I} \neq \emptyset$ **do**

$X := U$

$Y := Q^H X$

$Y := \text{MultishiftSolve}(G, z(\mathcal{I}), Y)$

$Y := QY$

$W := Q^H \text{sign}(Y)$

$W := \text{MultishiftSolve}(G^H, \bar{z}(\mathcal{I}), W)$

$W := QW$

$k := k + 1$

foreach $s = \mathcal{I}(t)$ **do**

$x := X(:, t)$, $y := Y(:, t)$, $w := W(:, t)$

if $(k \geq 2 \text{ and } \|w\|_\infty \leq w^H x)$ **or** $k \geq 5$ **then**

$f(s) = \|y\|_1$

 Delete $\mathcal{I}(t)$ and $U(:, t)$

else $U(:, t) := e_j$ where $|w(j)| = \|w\|_\infty$

...

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

...

// Take the maximum between the current estimates and a heuristic

$B := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1, s=0:m-1}$ // All columns are equal

$X := Q^H B$ // All columns are equal

$X := \text{MultishiftSolve}(G, z, X)$

$X := QX$

foreach $s = 0, \dots, m-1$ **do** $f(s) := \max(f(s), \|X(:, s)\|_1)$

Again: block algorithm from [Higham/Tisseur-2000] should be used in practice

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

...

// Take the maximum between the current estimates and a heuristic

$B := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1, s=0:m-1}$ // All columns are equal

$X := Q^H B$ // All columns are equal

$X := \text{MultishiftSolve}(G, z, X)$

$X := QX$

foreach $s = 0, \dots, m-1$ **do** $f(s) := \max(f(s), \|X(:, s)\|_1)$

Again: block algorithm from [Higham/Tisseur-2000] should be used in practice

Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...

Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...


Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...

FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 10 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$


30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 20 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$


30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 30 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$


30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 40 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$


30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)




FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 50 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$


30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)




FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 10 its




FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 20 its




FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 30 its



FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 40 its




FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 50 its




FoxLi(20k), $\Omega = (-1.2, 1.2)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




FoxLi(20k), $\Omega = (-0.12, 0.12)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



FoxLi(20k), $\Omega = (-0.012, 0.012)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



Uniform(15k), $\Omega = (-64, 64)^2$, 256^2 pix, 10 its

30 sec/iter on 256 cores of Stampede
(256 r.h.s./core, >4 TFlops)



Uniform(15k), $\Omega = (-64, 64)^2$, 256^2 pix, 20 its

30 sec/iter on 256 cores of Stampede
(256 r.h.s./core, >4 TFlops)



Uniform(15k), $\Omega = (-64, 64)^2$, 256^2 pix, 30 its

30 sec/iter on 256 cores of Stampede
(256 r.h.s./core, >4 TFlops)



Uniform(15k), $\Omega = (-64, 64)^2$, 256^2 pix, 40 its

30 sec/iter on 256 cores of Stampede
(256 r.h.s./core, >4 TFlops)




Uniform(15k), $\Omega = (-64, 64)^2$, 256^2 pix, 50 its

30 sec/iter on 256 cores of Stampede
(256 r.h.s./core, >4 TFlops)




Uniform(15k) work, $\Omega = (-64, 64)^2$, 256^2 pix, 10 its

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




Uniform(15k) work, $\Omega = (-64, 64)^2$, 256^2 pix, 20 its

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




Uniform(15k) work, $\Omega = (-64, 64)^2$, 256^2 pix, 30 its

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




Uniform(15k) work, $\Omega = (-64, 64)^2$, 256^2 pix, 40 its

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




Uniform(15k) work, $\Omega = (-64, 64)^2$, 256^2 pix, 50 its

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur




Uniform(20k), $\Omega = (-103, 103)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



Uniform(20k), $\Omega = (-10.3, 10.3)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



Future work

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ Extend EigTool to support high-performance interleaved algorithms (perhaps with accelerator support for the TRSM-like operation)
- ▶ More intelligent interpolation and stopping criteria
- ▶ Optional projection onto relevant eigenspaces

Future work

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ Extend EigTool to support high-performance interleaved algorithms (perhaps with accelerator support for the TRSM-like operation)
- ▶ More intelligent interpolation and stopping criteria
- ▶ Optional projection onto relevant eigenspaces

Future work

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ Extend EigTool to support high-performance interleaved algorithms (perhaps with accelerator support for the TRSM-like operation)
- ▶ More intelligent interpolation and stopping criteria
- ▶ Optional projection onto relevant eigenspaces

Future work

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ Extend EigTool to support high-performance interleaved algorithms (perhaps with accelerator support for the TRSM-like operation)
- ▶ More intelligent interpolation and stopping criteria
- ▶ Optional projection onto relevant eigenspaces

Acknowledgments and Availability

Support



Computational resources



Thanks:

Michael Saunders for extended discussions on regularizing QSD systems
Stephen Boyd and AJ Friend for extended discussions on IPMs

Minisymposium organizers:

Piotr Luszczek, Stan Tomov, and Azzam Haidar

Availability

Elemental is available under the New BSD License at libelemental.org
(Come see my poster on Elemental in the CSE Software mini on Monday!)

Questions?

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.