

High-performance computation of pseudospectra

Jack Poulson

Department of Mathematics
Stanford University

Innovative Computing Laboratory
The University of Tennessee
October 24, 2014

A note on collaborators

Multishift Hessenberg solves: collaboration with Greg Henry

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
 - ▶ High-performance batched analogues
 - ▶ A brief example of the python interface
 - ▶ Results
 - ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Overview of talk

- ▶ Van Loan's algorithm and the Demmel matrix
- ▶ Brief overview of pseudospectra
- ▶ Previous work
- ▶ High-performance batched analogues
- ▶ A brief example of the python interface
- ▶ Results
- ▶ Conclusions and future work

Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent over vertical line in complex plane [Van Loan-1985]
- ▶ Relationship to nearest stable matrix based on false conjecture, counter-example given in [Demmel-1987]
- ▶ Counter-example was first pseudospectral (computer) plot

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent over vertical line in complex plane [Van Loan-1985]
- ▶ Relationship to nearest stable matrix based on false conjecture, counter-example given in [Demmel-1987]
- ▶ Counter-example was first pseudospectral (computer) plot

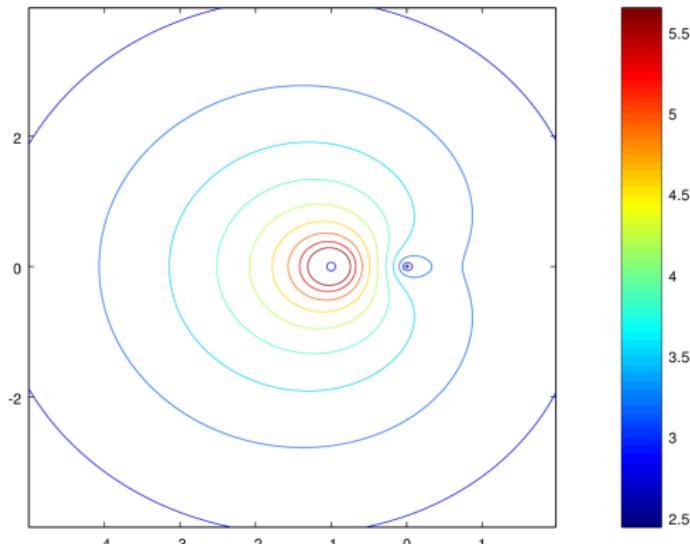
Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent over vertical line in complex plane [Van Loan-1985]
- ▶ Relationship to nearest stable matrix based on false conjecture, counter-example given in [Demmel-1987]
- ▶ Counter-example was first pseudospectral (computer) plot

Van Loan's algorithm and the Demmel matrix

- ▶ Efficient algorithm for evaluating resolvent over vertical line in complex plane [Van Loan-1985]
- ▶ Relationship to nearest stable matrix based on false conjecture, counter-example given in [Demmel-1987]
- ▶ Counter-example was first pseudospectral (computer) plot

$$D(n, \beta) = (\beta J_{-\beta^{-1}, n})^{-1}, \quad n = 3, \beta = 100$$



Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in [Trefethen and Embree's book *Spectra and pseudospectra*](#)...

Most common tool is EigTool [[Wright et al.-2001](#)], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in [Trefethen and Embree's book *Spectra and pseudospectra*](#)...

Most common tool is EigTool [[Wright et al.-2001](#)], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in Trefethen and Embree's book *Spectra and pseudospectra*...

Most common tool is EigTool [Wright et al.-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in **Trefethen and Embree's book *Spectra and pseudospectra***...

Most common tool is EigTool [Wright et al.-2001], but level 2 BLAS (trsv) and sequential

Pseudospectra

With convention $\|(\lambda I - A)^{-1}\|_p = \infty$ for $\lambda \in \mathcal{L}(A)$, eigenvalues are singularities of resolvent norm $\|(\xi I - A)^{-1}\|_p$.

Natural generalization of spectrum for each p -norm and $\epsilon > 0$:

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - A)^{-1}\|_p > \frac{1}{\epsilon} \right\}$$

Reinvented many times [Varah-1967, Landau-1975, Godunov et al.-1982, Trefethen-1990, Hinrichsen/Pritchard-1992,...]

Extensive review of field in [Trefethen and Embree's book *Spectra and pseudospectra*](#)...

Most common tool is EigTool [[Wright et al.-2001](#)], but level 2 BLAS (trsv) and sequential

Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

Relationship to previous work

- ▶ “Embarrassingly parallel” Lanczos [Braconnier-1996]
- ▶ Parallel sparse-direct shift-and-invert [Fraysse et al.-1996]
- ▶ Parallel path-following methods, e.g., [Mezher-2001,Bekas et al.-2000,2001,...]
- ▶ Our focus on level 3 (parallel) batch evaluation over point clouds after parallel Schur decomposition
- ▶ Key idea: simultaneously drive many Lanczos/Arnoldi methods via multishift extensions of standard TRSM algorithm
- ▶ Software for one and two-norm pseudospectra already released in Elemental [P. et al.-2013]

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - G)^{-1}\|_2 > \frac{1}{\epsilon} \right\},$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - G)^{-1}\|_2 > \frac{1}{\epsilon} \right\},$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm

[Van Loan-1985,Lui-1997]

Given a reduction to condensed form, $A = QGQ^H$,

$$\mathcal{L}_\epsilon^p(A) = \left\{ \xi \in \mathbb{C} : \|Q(\xi I - G)^{-1}Q^H\|_p > \frac{1}{\epsilon} \right\},$$

and, when $p = 2$,

$$\mathcal{L}_\epsilon^2(A) = \left\{ \xi \in \mathbb{C} : \|(\xi I - G)^{-1}\|_2 > \frac{1}{\epsilon} \right\},$$

First reduce to (quasi-)triangular/Hessenberg form, then,

- ▶ (Restarted) Lanczos/Arnoldi for each $\|(\xi I - G)^{-1}\|_2$, or
- ▶ Blocked Hager [Higham/Tisseur-2000] for each $\|Q(\xi I - G)^{-1}Q^H\|_1$.

Note that this algorithm makes less sense when $(\xi I - A)^{-1}$ can already be applied in quadratic time (e.g., if A is 3D FEM discretization)

(Extended) Van Loan algorithm for $\mathcal{L}_\epsilon^2(A)$

Algorithm: Two-norm pseudospectra via extended Van Loan algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shifts $\Omega \subset \mathbb{C}$, restart size k

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(\xi I - A)^{-1}\|_2\}_{\xi \in \Omega}$

$G := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

// Estimate $\|(\xi I - G)^{-1}\|_2$ via Restarted Arnoldi

Choose $v_0 \in \mathbb{C}^n$ with $\|v_0\|_2 = 1$

while not converged **do**

for $j = 0, \dots, k - 1$ **do**

// $(\xi I - G)^{-H}(\xi I - G)^{-1} V_j = V_j H_j + v_j (\beta_j e_j)^H$

$x_j := (\xi I - G)^{-H}(\xi I - G)^{-1} v_j$

Expand Arnoldi decomposition using x_j

$[\lambda, v_0] := \text{MaxEig}(H_k)$

$\phi(\xi) := \text{RealPart}(\lambda)$

[Hager-1984,Higham-1988] approach for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, $\Omega \subset \mathbb{C}$

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(A - \xi I)^{-1}\|_1\}_{\xi \in \Omega}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

// Estimate $\|(A - \xi I)^{-1}\|_1$ via Hager-Higham algorithm

Choose $u \in \mathbb{C}^n$ with $\|u\|_1 = 1$

$k := 0$

repeat

$x := u$

$y := Q(G - \xi I)^{-1}Q^H x$

$w := Q(G - \xi I)^{-H}Q^H \text{sign}(y)$

$u := e_j$ where $|w(j)| = \|w\|_\infty$

$k := k + 1$

until $(k \geq 2 \text{ and } \|w\|_\infty \leq w^H x) \text{ or } k \geq 5$

$\phi(\xi) := \|y\|_1$

// Take the maximum between the current estimate and a heuristic

$b := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1}$

$x := Q(G - \xi I)^{-1}Q^H b$

$\phi(\xi) := \max(\phi(\xi), \|x\|_1)$

Block algorithm [Higham/Tisseur-2000] should be used in practice

[Hager-1984,Higham-1988] approach for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, $\Omega \subset \mathbb{C}$

Output: $\{\phi(\xi)\}_{\xi \in \Omega} \approx \{\|(A - \xi I)^{-1}\|_1\}_{\xi \in \Omega}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

foreach $\xi \in \Omega$ **do**

// Estimate $\|(A - \xi I)^{-1}\|_1$ via Hager-Higham algorithm

Choose $u \in \mathbb{C}^n$ with $\|u\|_1 = 1$

$k := 0$

repeat

$x := u$

$y := Q(G - \xi I)^{-1}Q^H x$

$w := Q(G - \xi I)^{-H}Q^H \text{sign}(y)$

$u := e_j$ where $|w(j)| = \|w\|_\infty$

$k := k + 1$

until $(k \geq 2 \text{ and } \|w\|_\infty \leq w^H x) \text{ or } k \geq 5$

$\phi(\xi) := \|y\|_1$

// Take the maximum between the current estimate and a heuristic

$b := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1}$

$x := Q(G - \xi I)^{-1}Q^H b$

$\phi(\xi) := \max(\phi(\xi), \|x\|_1)$

Block algorithm [Higham/Tisseur-2000] should be used in practice

Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

(Generalized) multishift solves

Computing $\{(\delta_j F - G)^{-1} y_j\}_j$ equivalent to solving for X in

$$FXD - GX = Y,$$

where $D = \text{diag}((\delta_0, \delta_1, \dots))$.

- ▶ Multishift triangular solves (G triangular, $F = I$) require trivial changes to usual high-performance TRSM algorithms [Henry-1994]
- ▶ Quasi-triangular similar; adjust blocksizes to not split 2×2
- ▶ Very recently used by [Gates/Haidar/Dongarra-2014] for eigenvectors of nonsymmetric matrices
- ▶ Generalized multishift (quasi-)triangular solves can similarly be made high-performance with a simple trick...
- ▶ Multishift Hessenberg solves heavily investigated for control theory [Datta et al-1994, Henry-1994]
- ▶ Large fraction of work in Hessenberg case is level 1...

Interleaved Van Loan algorithm for $\mathcal{L}_\epsilon^2(A)$

Algorithm: Two-norm pseudospectra via interleaved extended Van Loan algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$, restart size k

Output: $f \approx [\|(z(s)I - A)^{-1}\|_2]_{s=0:m-1}$

$G := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

Initialize each column of $W_0 \in \mathbb{C}^{n \times m}$ to have unit two-norm

$\mathcal{I} := (0, \dots, m-1)$

while $\mathcal{I} \neq \emptyset$ **do**

for $j = 0, \dots, k-1$ **do**

// $(z(s)I - G)^{-H}(z(s)I - G)^{-1}W_j(t) = W_j(t)\mathcal{H}_j(t) + W_j(:, t)(b_j(t)e_j)^H, \forall s = \mathcal{I}(t)$

$X := \text{MultishiftSolve}(G, z(\mathcal{I}), W_j)$

$X := \text{MultishiftSolve}(G^H, \bar{z}(\mathcal{I}), X)$

Expand Arnoldi decompositions

foreach $s = \mathcal{I}(t)$ **do**

$[\lambda, W_0(:, t)] := \text{MaxEig}(\mathcal{H}_k(t))$

$f(s) := \text{RealPart}(\lambda)$

if converged **then**

Delete $\mathcal{I}(t)$ and $W_0(:, t)$

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

$[Q, G] := \text{Schur}(A)$, $\text{RealSchur}(A)$, or $\text{Hessenberg}(A)$

Initialize each column of $U \in \mathbb{C}^{n \times m}$ to have unit one-norm

$\mathcal{I} := (0, \dots, m-1)$

$k := 0$

while $\mathcal{I} \neq \emptyset$ **do**

$X := U$

$Y := Q^H X$

$Y := \text{MultishiftSolve}(G, z(\mathcal{I}), Y)$

$Y := QY$

$W := Q^H \text{sign}(Y)$

$W := \text{MultishiftSolve}(G^H, \bar{z}(\mathcal{I}), W)$

$W := QW$

$k := k + 1$

foreach $s = \mathcal{I}(t)$ **do**

$x := X(:, t)$, $y := Y(:, t)$, $w := W(:, t)$

if $(k \geq 2 \text{ and } \|w\|_\infty \leq w^H x) \text{ or } k \geq 5$ **then**

$f(s) = \|y\|_1$

Delete $\mathcal{I}(t)$ and $U(:, t)$

else $U(:, t) := e_j$ where $|w(j)| = \|w\|_\infty$

...

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

...

// Take the maximum between the current estimates and a heuristic

$B := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1,s=0:m-1}$ // All columns are equal

$X := Q^H B$ // All columns are equal

$X := \text{MultishiftSolve}(G, z, X)$

$X := QX$

foreach $s = 0, \dots, m-1$ **do** $f(s) := \max(f(s), \|X(:, s)\|_1)$

Again: block algorithm from [Higham/Tisseur-2000] should be used in practice

Interleaved Hager-Higham algorithm for $\mathcal{L}_\epsilon^1(A)$

Algorithm: One-norm pseudospectra via interleaved Hager-Higham algorithm

Input: $A \in \mathbb{F}^{n \times n}$, shift vector $z \in \mathbb{C}^m$

Output: $f \approx [\|(z(j)I - A)^{-1}\|_1]_{j=0:m-1}$

...

// Take the maximum between the current estimates and a heuristic

$B := \frac{2}{3n}[(-1)^j(n+j-1)]_{j=0:n-1,s=0:m-1}$ // All columns are equal

$X := Q^H B$ // All columns are equal

$X := \text{MultishiftSolve}(G, z, X)$

$X := QX$

foreach $s = 0, \dots, m-1$ **do** $f(s) := \max(f(s), \|X(:, s)\|_1)$

Again: block algorithm from [Higham/Tisseur-2000] should be used in practice

Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...

Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...

Batching and deflation

- ▶ Maximum number of simultaneous shifts constrained by memory
- ▶ Could bring in new shift after each deflation, but easier to break into batches
- ▶ In practice, only small number of iterations needed per shift...

Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

Investigating the Fox-Li matrix with python

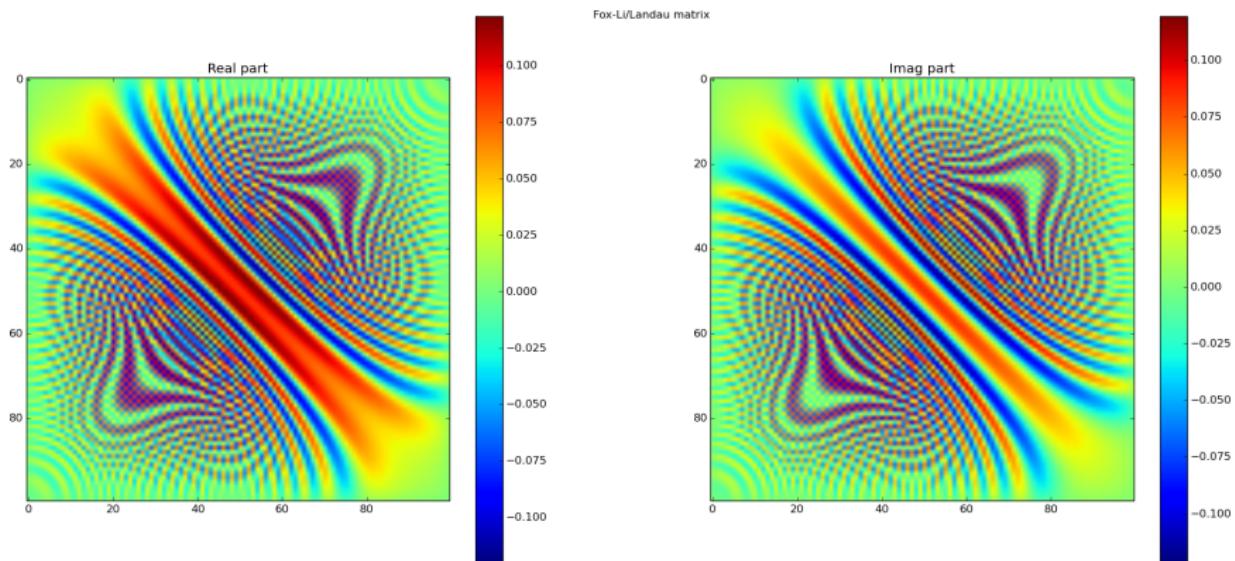
```
import math, El
n = 100                      # matrix size
realRes = imagRes = 100 # grid resolution

# Display an instance of the Fox–Li/Landau matrix
A = El.DistMatrix(El.zTag)
El.FoxLi(A,n)
El.Display(A,"Fox–Li matrix")

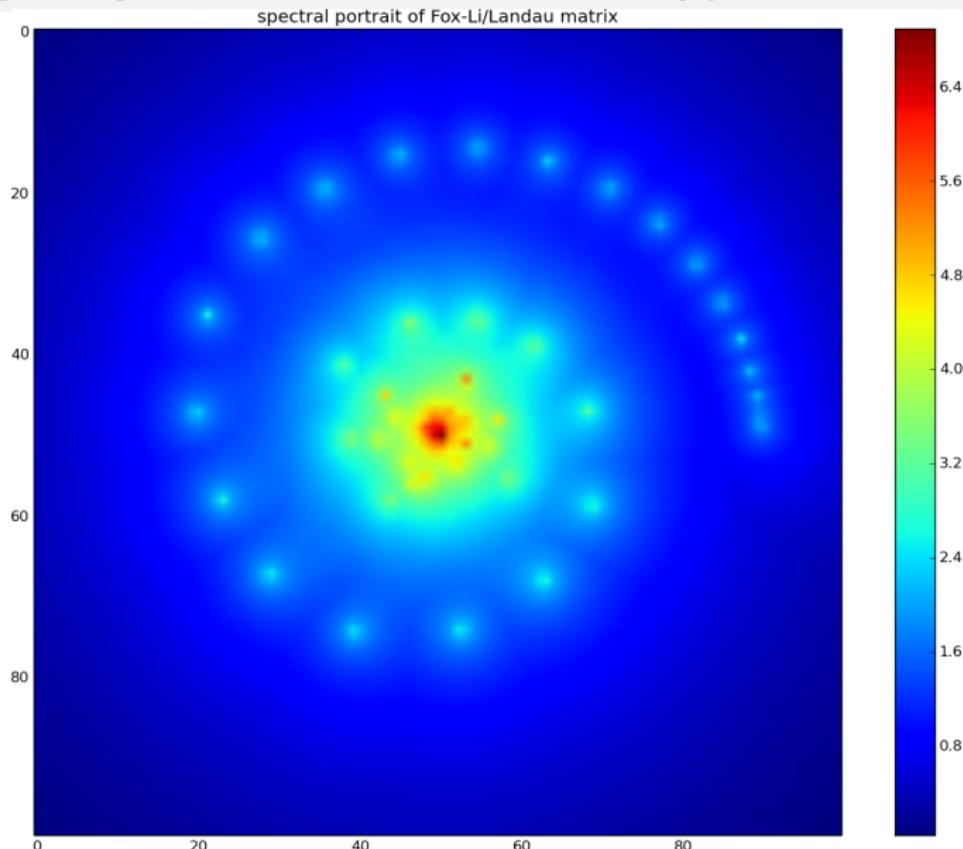
# Display its spectral portrait
portrait = El.SpectralPortrait(A,realRes,imagRes)
El.EntrywiseMap(portrait,math.log10)
El.Display(portrait,"Spectral portrait of Fox–Li matrix")

# Display its singular values
s = El.SVD(A)
El.EntrywiseMap(s,math.log10)
El.Display(s,"log10(svd(A))")
```

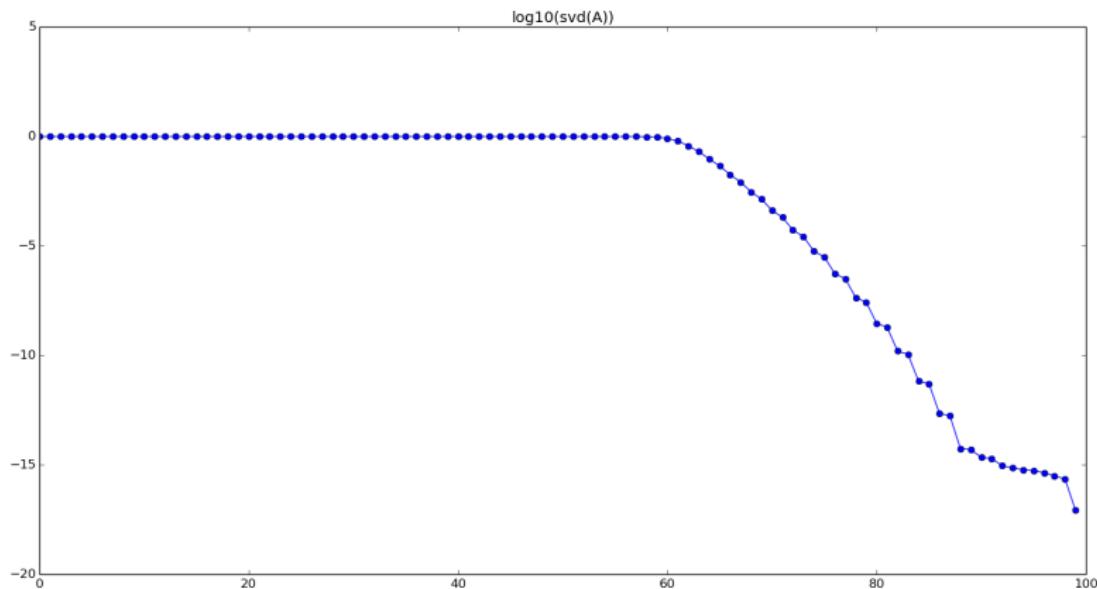
Investigating the Fox-Li matrix with python



Investigating the Fox-Li matrix with python



Investigating the Fox-Li matrix with python



Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

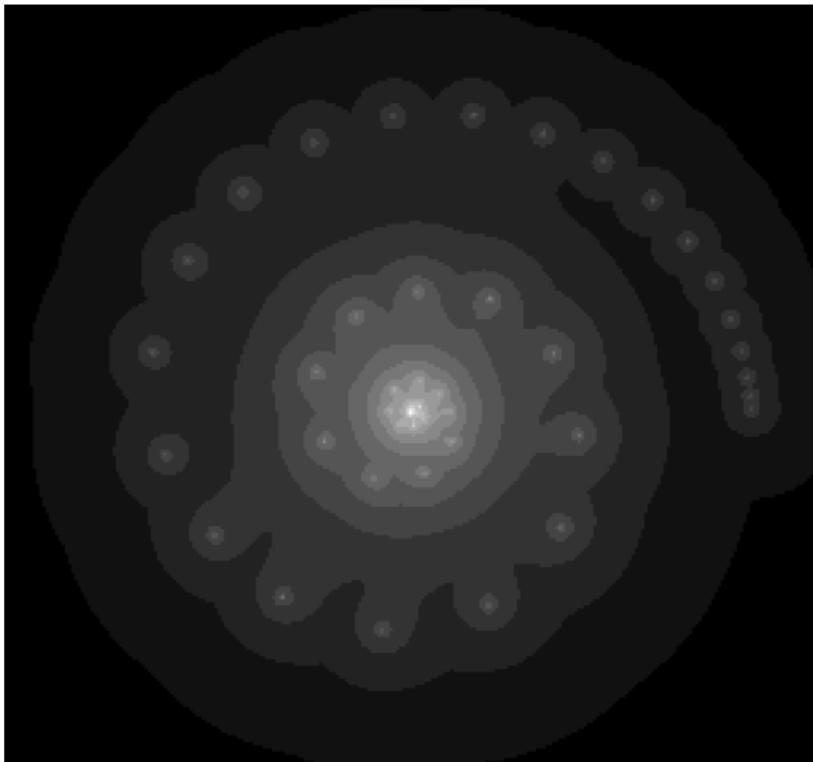
Results

Conclusions and future work

FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 10 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$

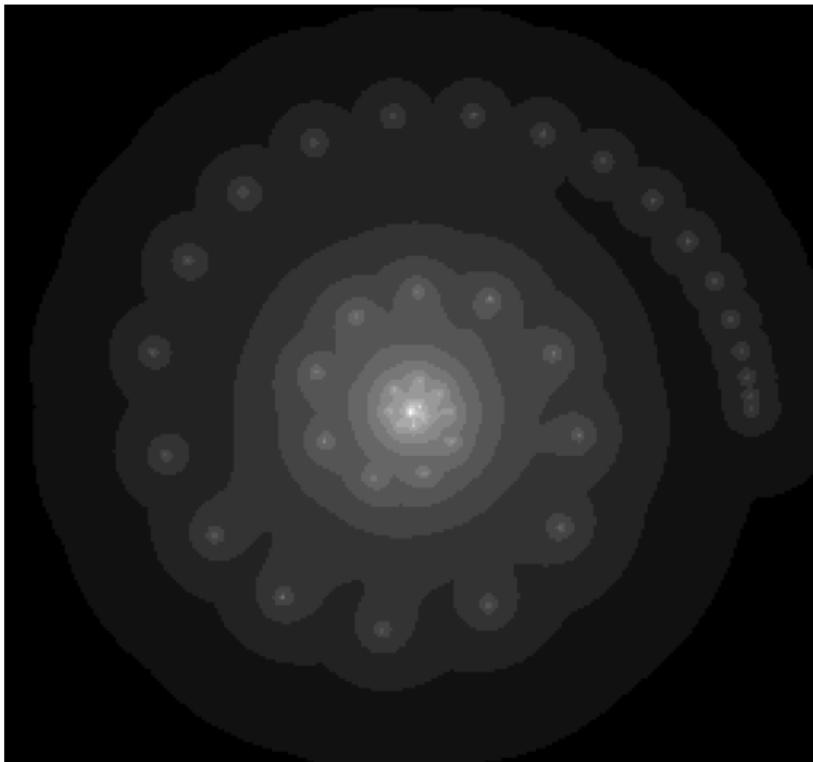
30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 20 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$

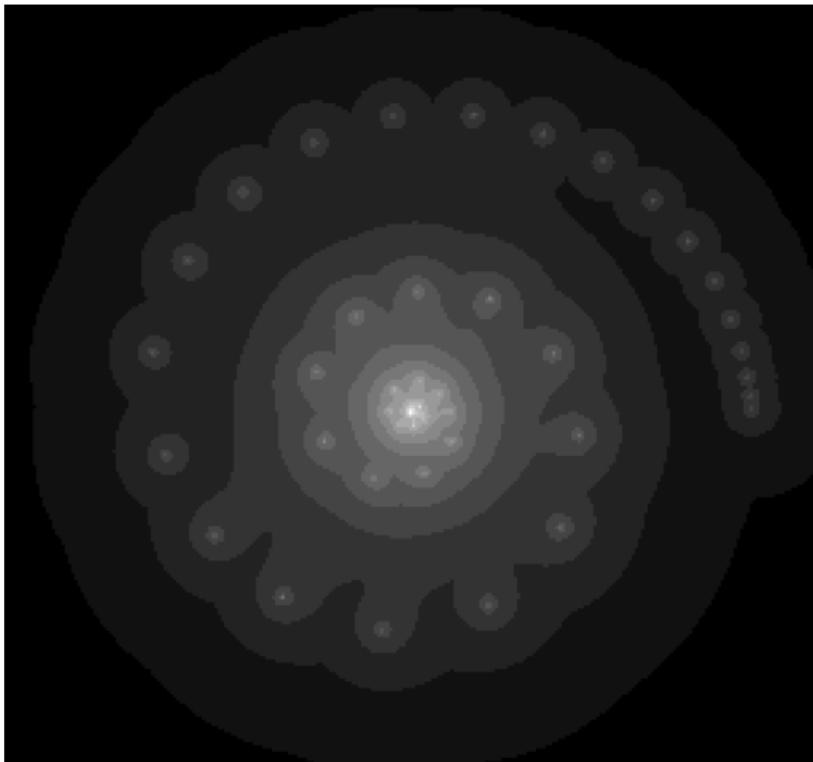
30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 30 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$

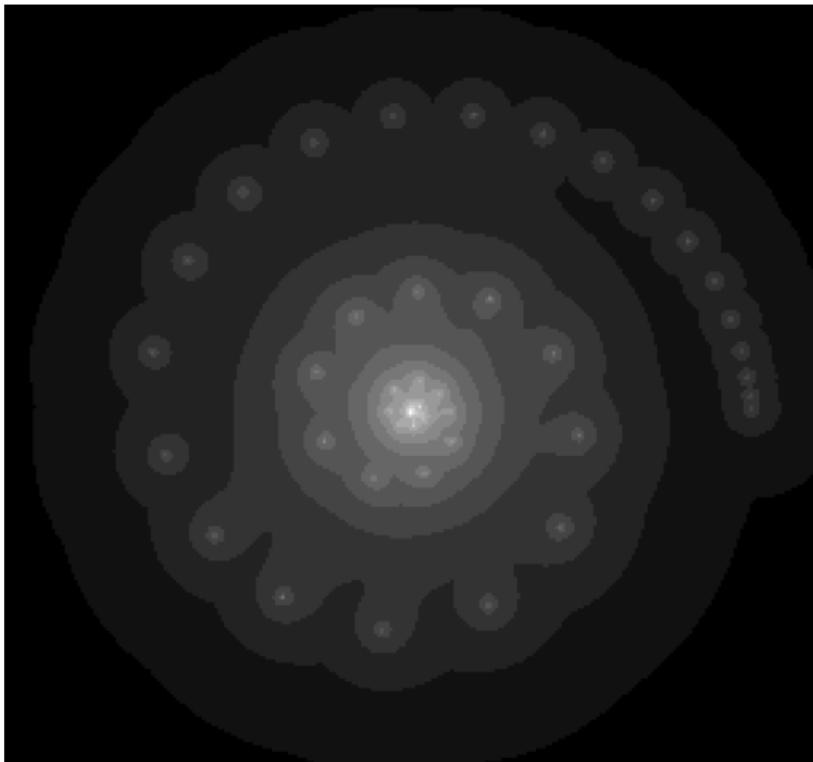
30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 40 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$

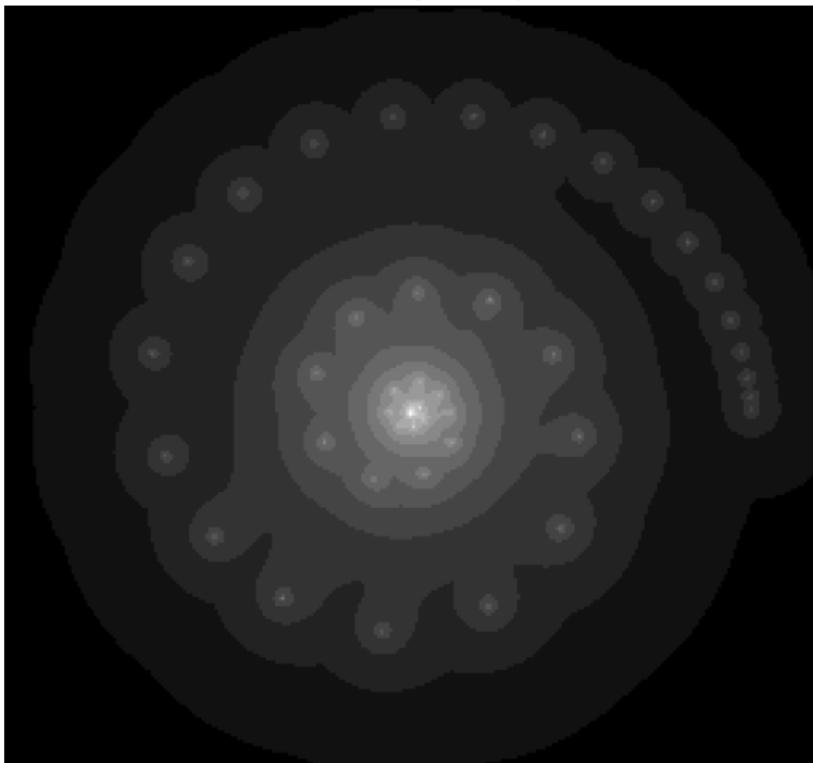
30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



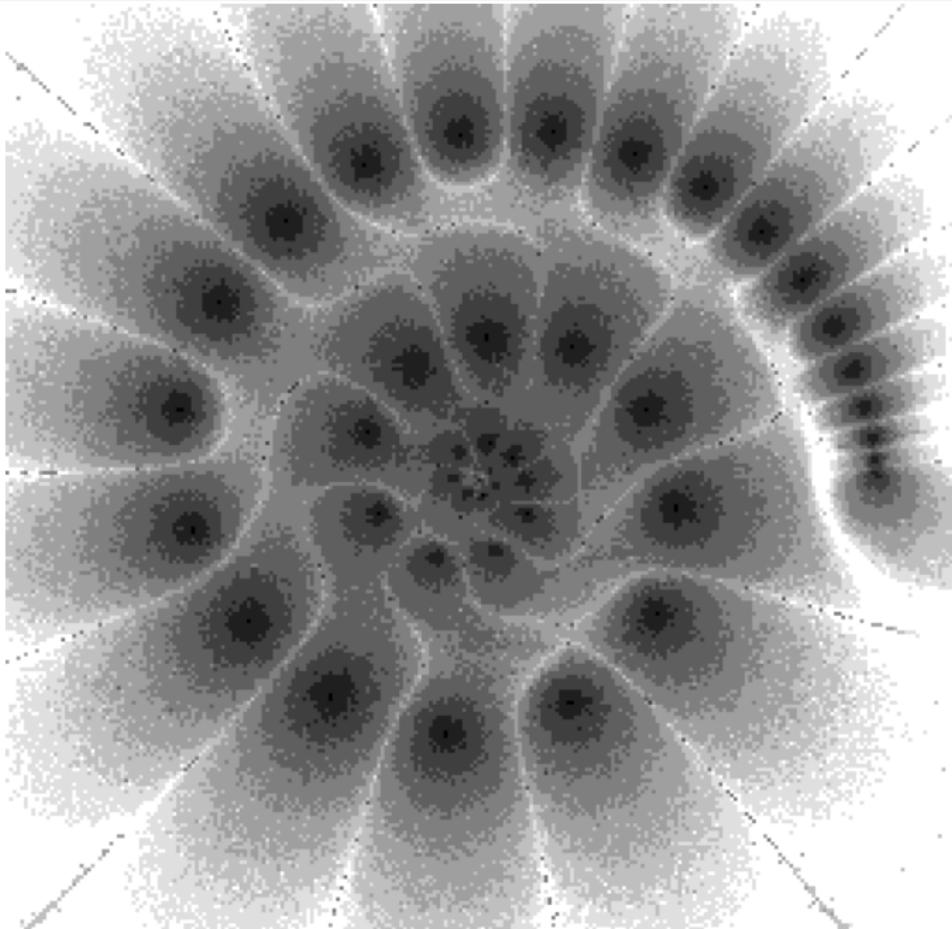
FoxLi(15k), $\Omega = (-1.2, 1.2)^2$, 256^2 pixels, 50 its

$$(\mathcal{A}u)(x) = \sqrt{iF/\pi} \int_{-1}^1 e^{iF(x-y)^2} u(y) dy, \quad F = 16\pi$$

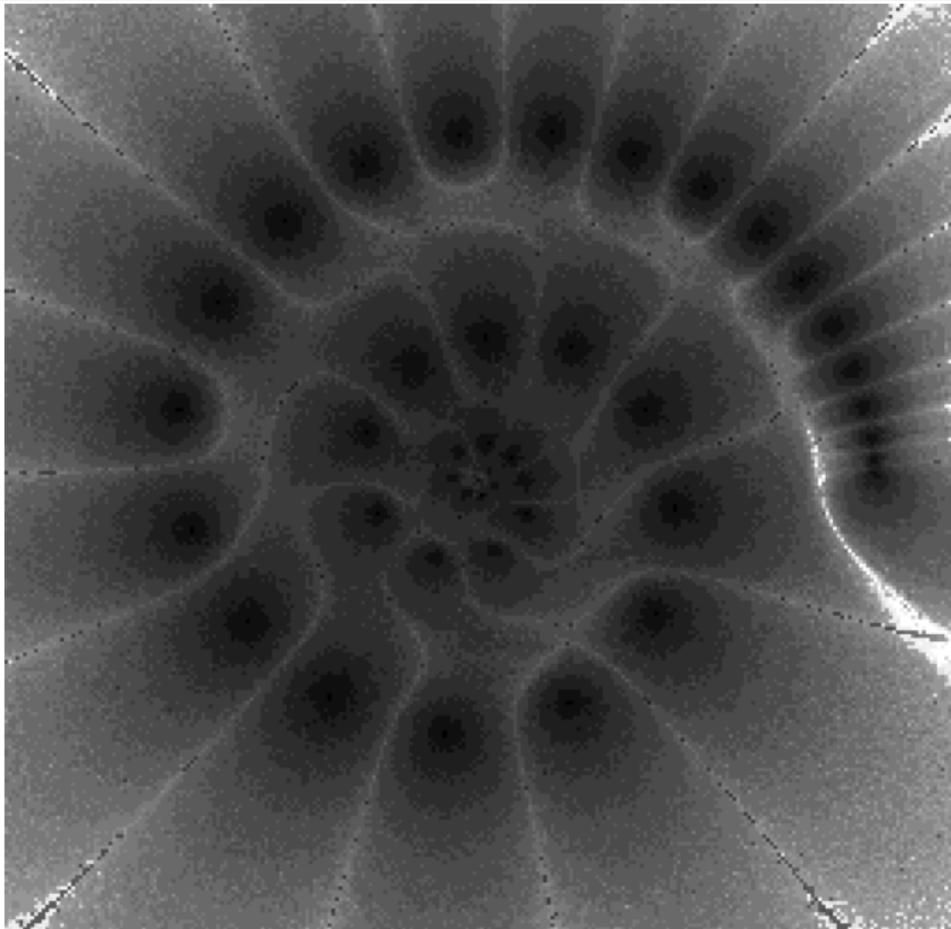
30 sec/iter on 256 cores of Stampede (256 r.h.s./core, >4 TFlops)



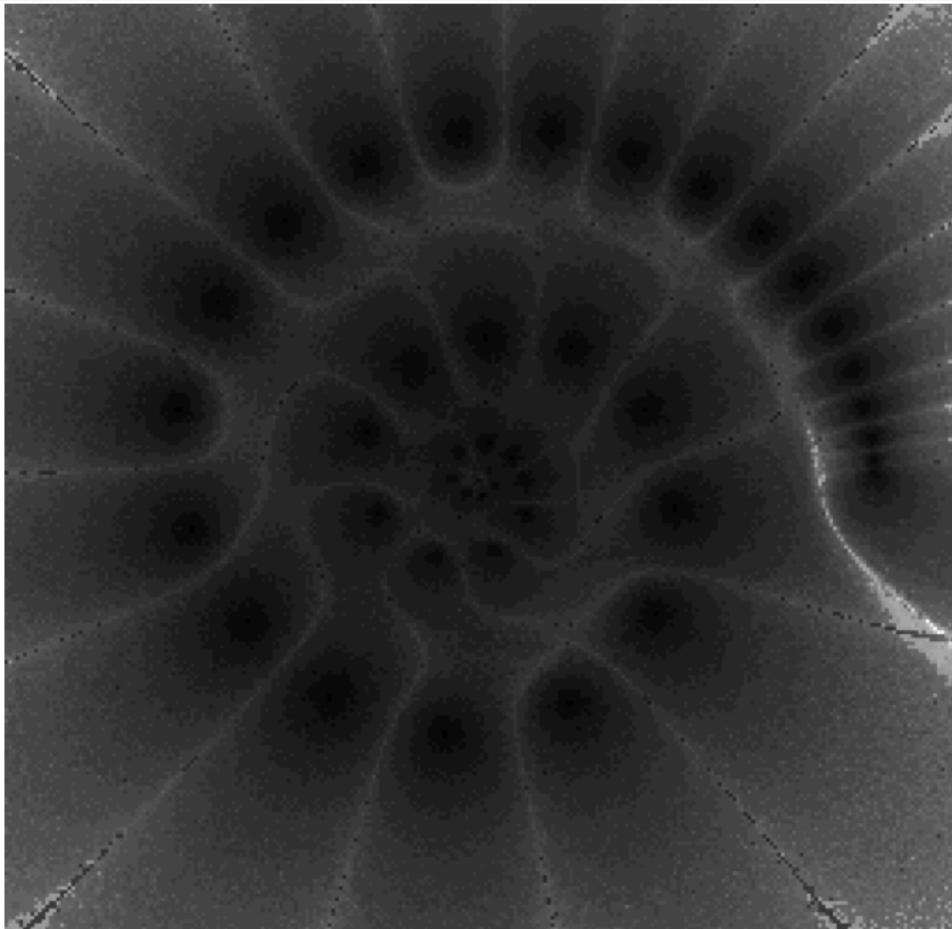
FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 10 its



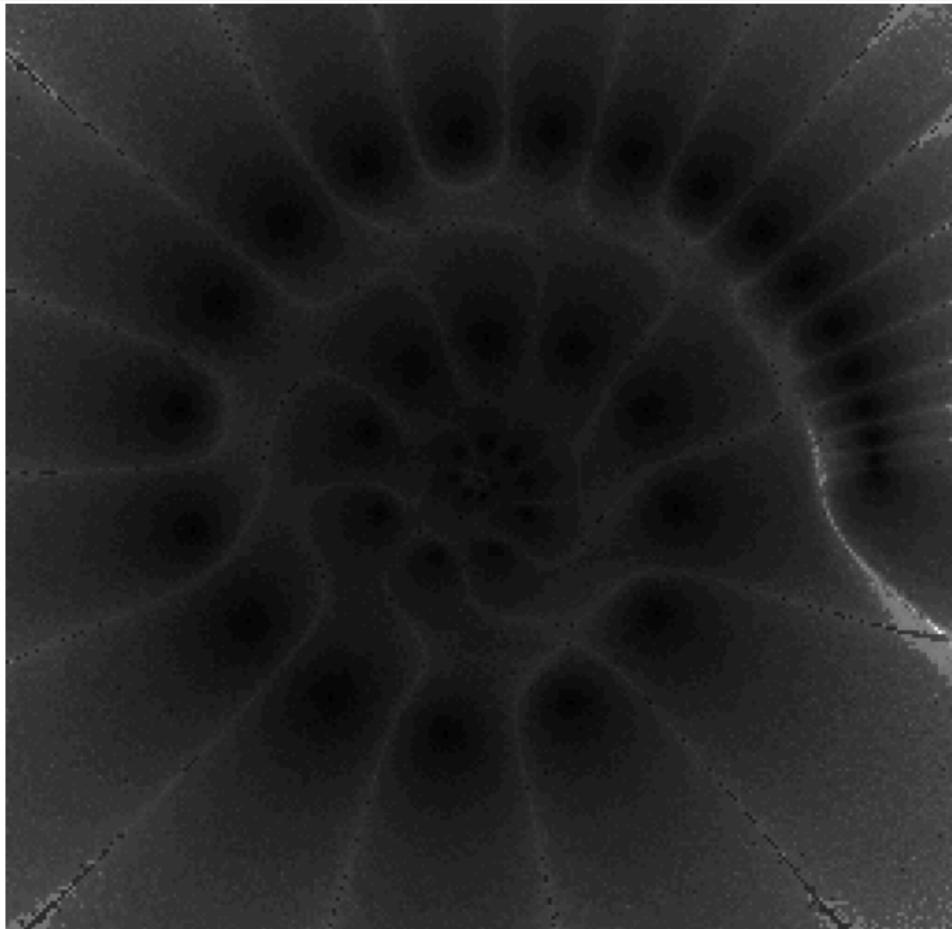
FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 20 its



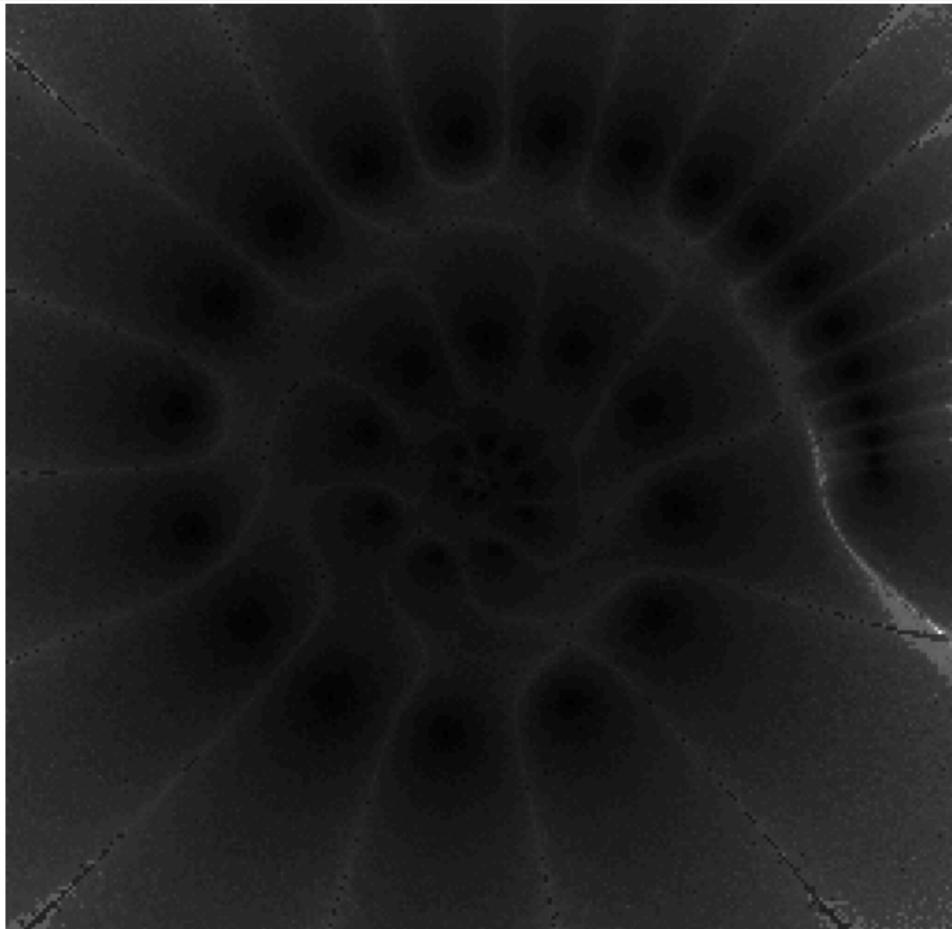
FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 30 its



FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 40 its

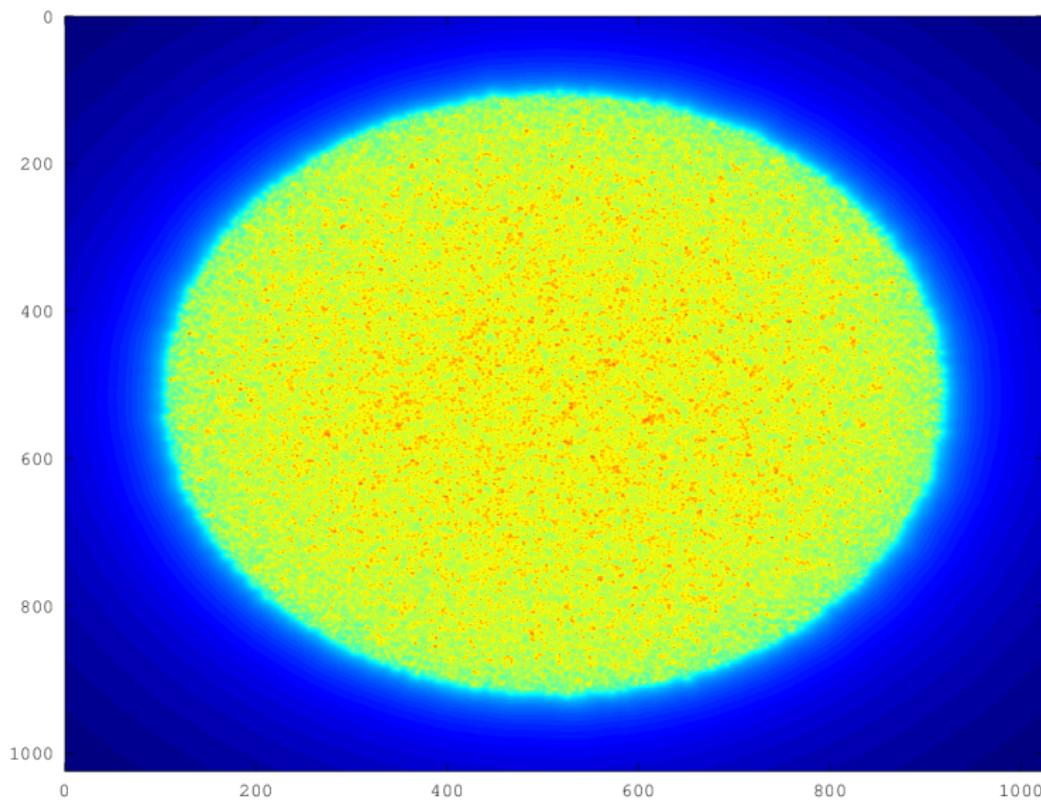


FoxLi(15k) work, $\Omega = (-1.2, 1.2)^2$, 256^2 pix, 50 its



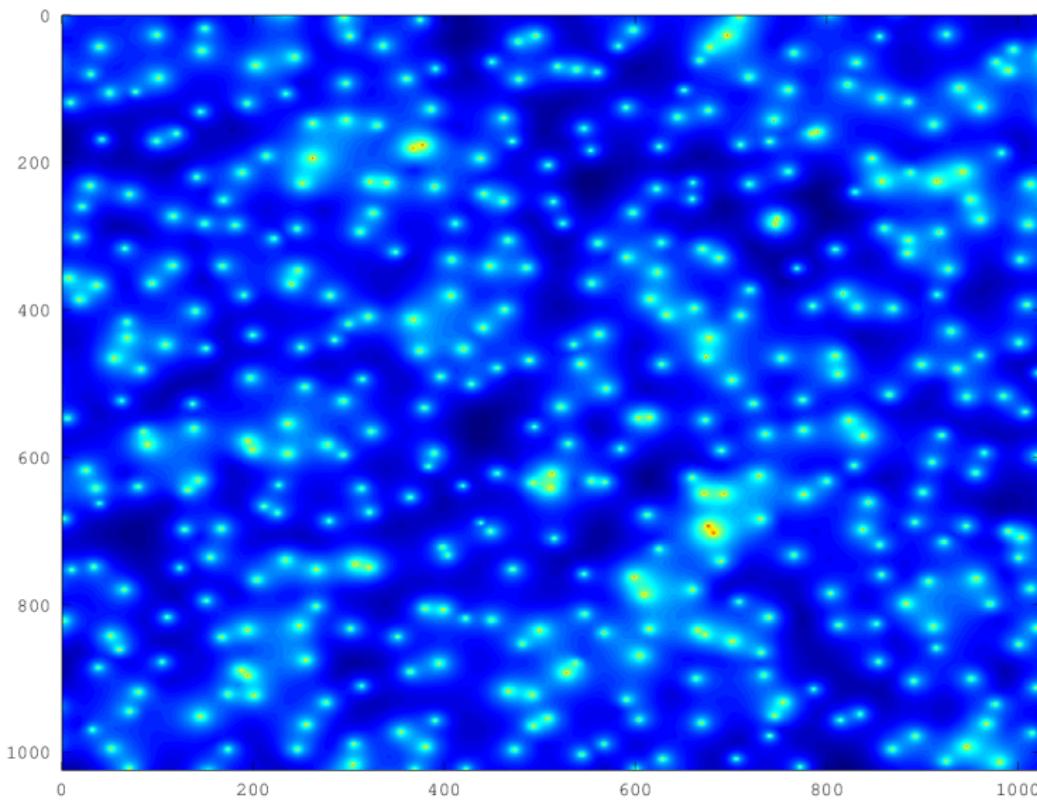
Uniform(20k), $\Omega = (-103, 103)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



Uniform(20k), $\Omega = (-10.3, 10.3)^2$, 1024^2 pix

4 256^2 pieces: 75 sec/iter on 256 nodes of Blue Gene/Q,
(64 r.h.s./core, >11 TFlops), 1175 sec for Schur



Outline

Van Loan's algorithm and the Demmel matrix

Brief overview of pseudospectra

Previous work

High-performance batched analogues

A brief example of the python interface

Results

Conclusions and future work

Convergence issues

- ▶ Typically small number of iterations required for visual convergence.
- ▶ Given current batch iteration, can output image after each fixed number of iterations
- ▶ Problematic pixels typically uninteresting (resolvent gradient is small)

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Future directions

- ▶ Interleaved block one-norm pseudospectral algorithm
- ▶ More intelligent interpolation and stopping criteria
- ▶ Investigate (preconditioned) 2D/3D wave equations
- ▶ Better understanding of practical limits of SDC EVD
- ▶ Optional projection onto relevant eigenspaces
- ▶ Complex distributed Hessenberg QR/QZ with AED...
- ▶ Black-box high-performance generalized pseudospectra

Acknowledgments and Availability

Support



Computational resources



My hosts

Jakub Kurzak and Jack Dongarra

Availability

Elemental is available under the New BSD License at
libelemental.org

Questions?

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.

Generalized multishift (quasi-)TRSM

$$FXD - GX = Y$$

As long as a 2x2 block is not split:

$$\begin{pmatrix} (F_{1,1}X_1D - G_{1,1}X_1) + (F_{1,2}X_2D - G_{1,2}X_2) \\ F_{2,2}X_2D - G_{2,2}X_2 \end{pmatrix} = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix}$$

If $F_{2,2}$ and $G_{2,2}$ are small and square, have each process locally solve for a few right-hand sides of X_2 . Then, form

$$\hat{Y}_1 := Y_1 - (F_{1,2}X_2D - G_{1,2}X_2)$$

via a parallel GEMM with each column of $F_{1,2}X_2$ appropriately scaled afterwards.

All that is left is to recurse on

$$F_{1,1}X_1D - G_{1,1}X_1 = \hat{Y}_1.$$

When $F = I$, algorithm is much simpler.